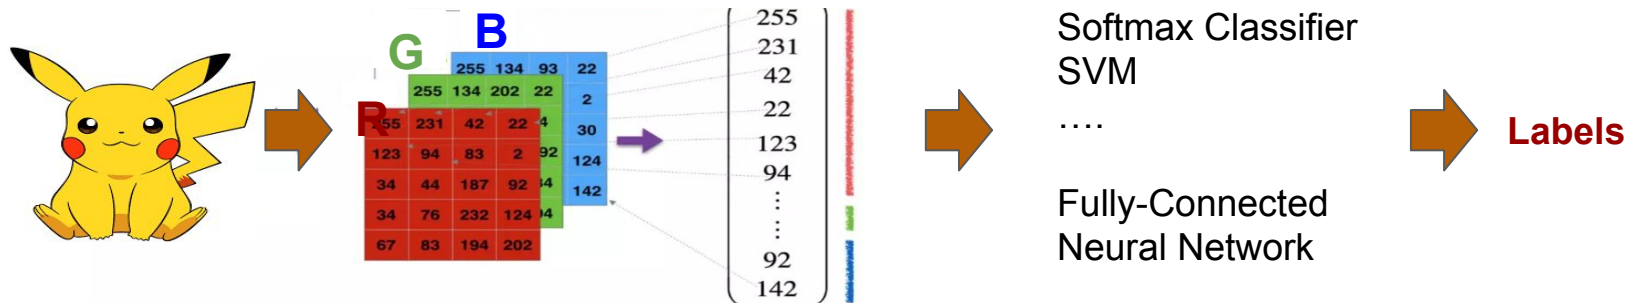



Convolutional Neural Network

Before CNN

Computers See Image



Think about MNIST Dataset



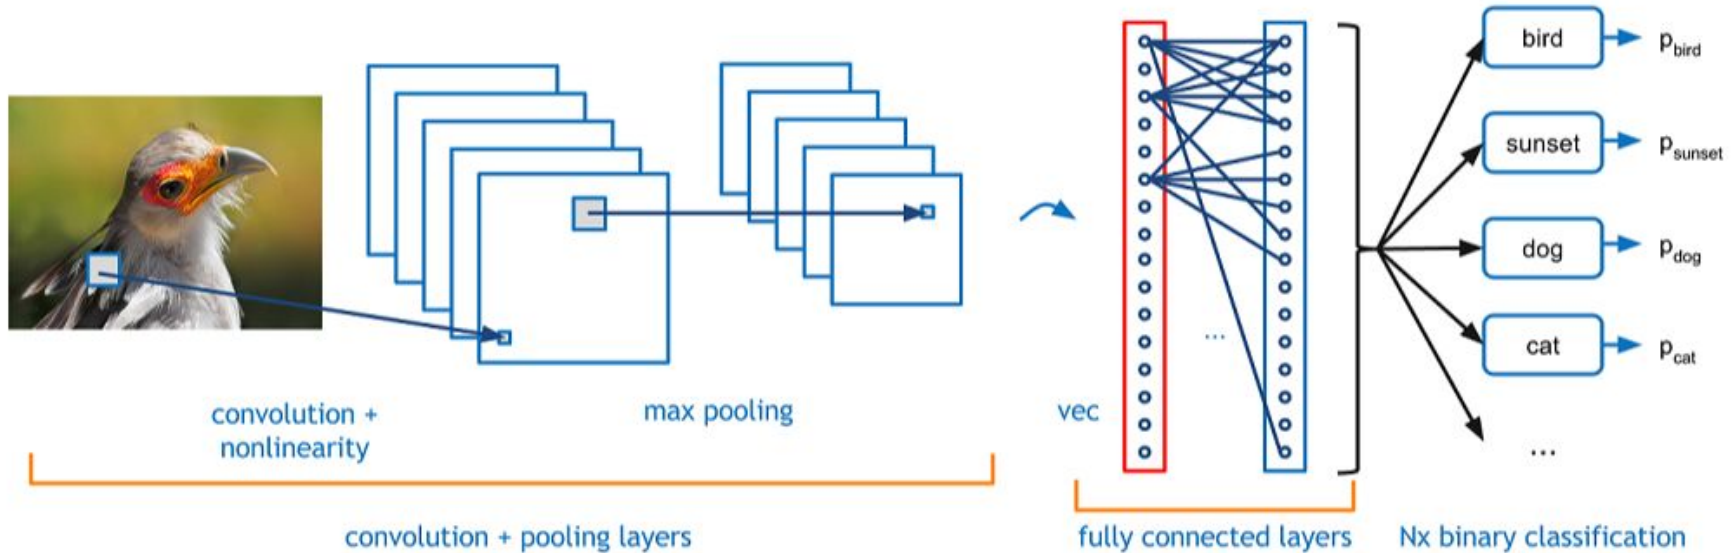
The above model requires the digit should be in the center of the image and it had to be the only thing in the image.

Intro to CNN



https://www.youtube.com/watch?v=FwFduRA_L6Q

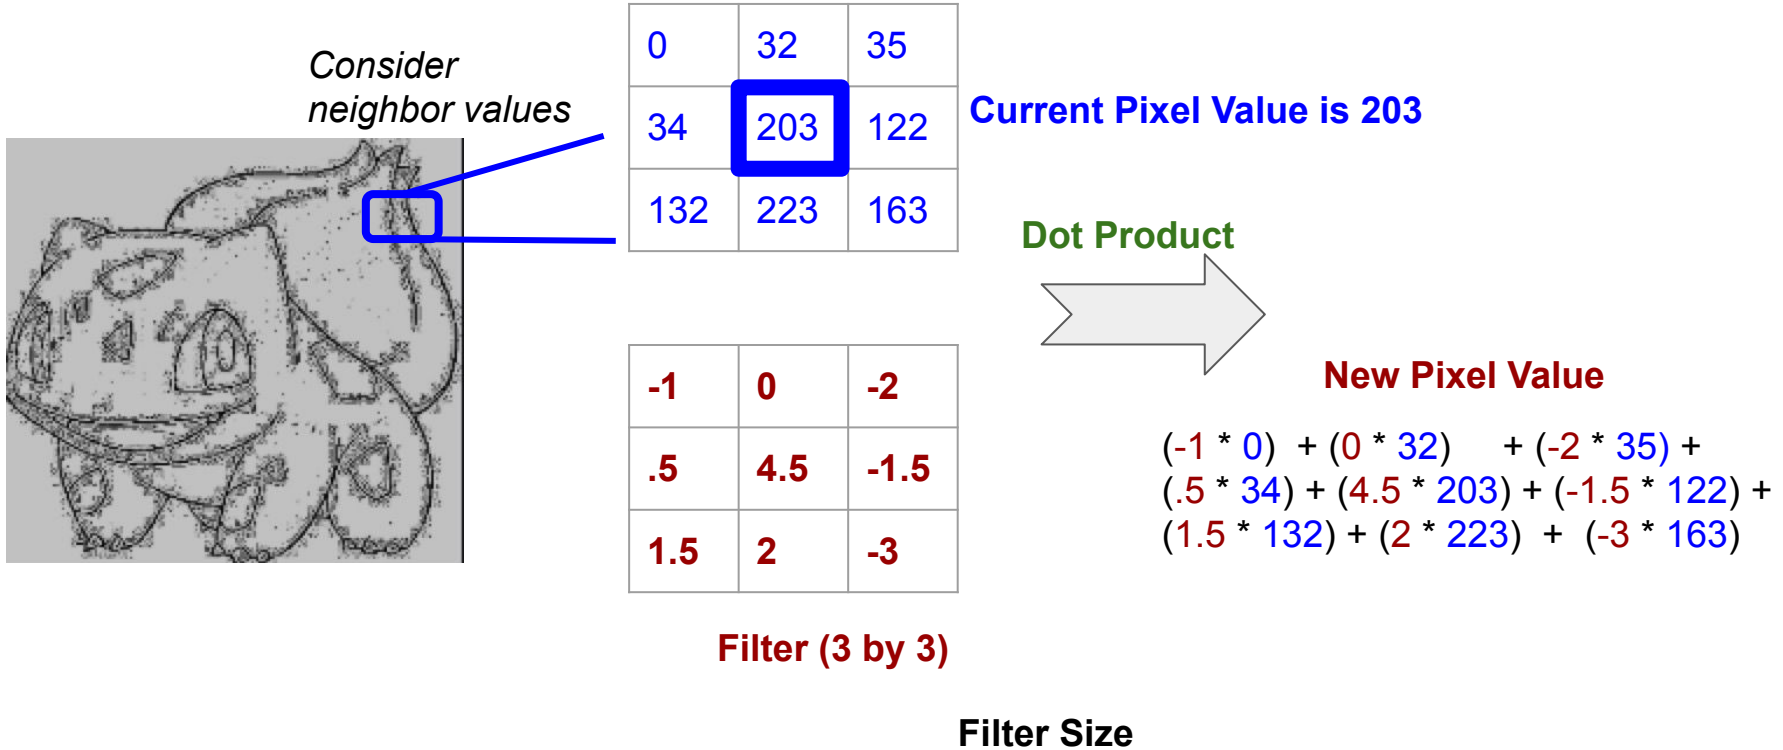
Convolutional Neural Network



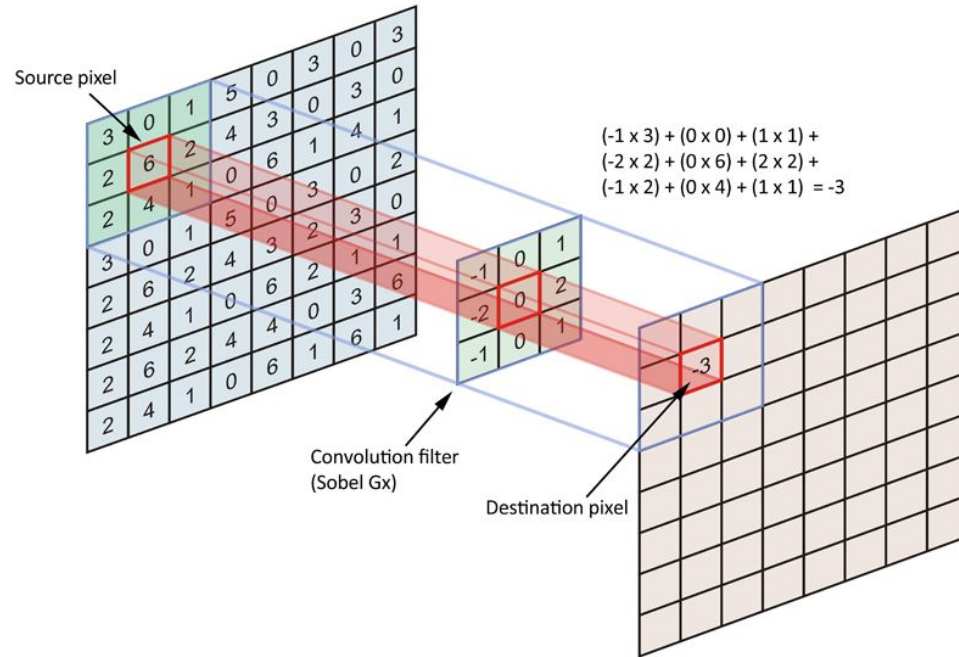
**Extracting useful
features of data**

**Perform a ML task (like
classification based on the
vectorized data)**

Filter Operation



Filter Operation



The intent of convolution is to encode source data matrix (entire image) in terms of a filter or kernel. More specifically, we are trying to encode the pixels in the **neighborhood** of **anchor/source** pixels

<https://datascience.stackexchange.com/questions/23183/why-convolutions-always-use-odd-number-s-as-filter-size>

Convolutional Operation

- Apply the **same** filter for every pixel in the original image
- Filter Size is the shape of the filter matrix (yellow one)

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image



4		

Convolved
Feature

*Feature
Map*

Stride Size

- Controls how the filter move around the image
- It is the amount by which the filter shifts



Padding Size

- Pads the image with zeros around the **border**
- Make the input image and feature map have the same spatial dimensions

0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

Kernel

0	-1	0
-1	5	-1
0	-1	0

114				

Stride: 1

Size of zero padding:

$$(k-1)/2$$

<https://stackoverflow.com/questions/52067833/how-to-plot-an-animated-matrix-in-matplotlib>

Convolutional Operation

- Filter Size: K
- Stride Size: S
- Padding Size: P

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Input size

$$O = \frac{W - K + 2P}{S} + 1$$

Output size

Multi-Channel CNN

- A color image is a 3-D tensor
- 400 (height) 630 (width) 3 (R,G,B channels)

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

308

+

1	0	0
0	1	-1
1	0	-1

Kernel Channel #2

-498

+

0	1	1
0	1	0
1	-1	1

Kernel Channel #3

164

+ 1 = -25
Bias = 1


-25			...
			...
			...
			...
...

Output

```
from matplotlib.image import imread
import numpy as np
img = imread('pikka_3.jpg')

print(img.shape)

plt.imshow(img, interpolation='nearest')
<matplotlib.image.AxesImage at 0x11b404278>
```



From Keras Layers Conv2D

Input shape


4D tensor with shape: (batch, channels, rows, cols) if data_format is "channels_first" or 4D tensor with shape: (batch, rows, cols, channels) if data_format is "channels_last".

Output shape

4D tensor with shape: (batch, filters, new_rows, new_cols) if data_format is "channels_first" or 4D tensor with shape: (batch, new_rows, new_cols, filters) if data_format is "channels_last". rows and cols values might have changed due to padding.

https://www.researchgate.net/post/How_will_channels_RGB_effect_convolutional_neural_network

Filter comes from “Image Processing”

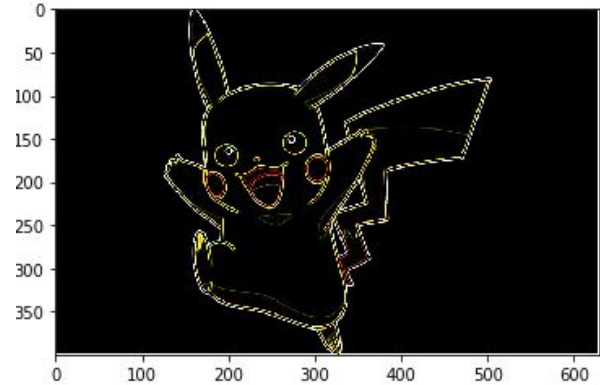


Image

```
print(kernel)
```

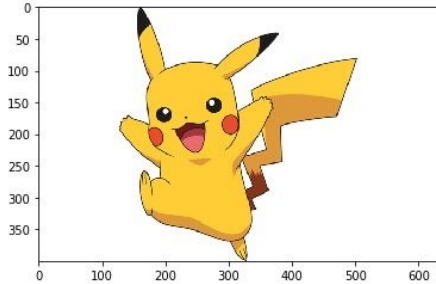
```
[[ -1  -1  -1]
 [ -1   8  -1]
 [ -1  -1  -1]]
```

**Edge
Detection**



**Convolved
Features**

Filter comes from “Image Processing”

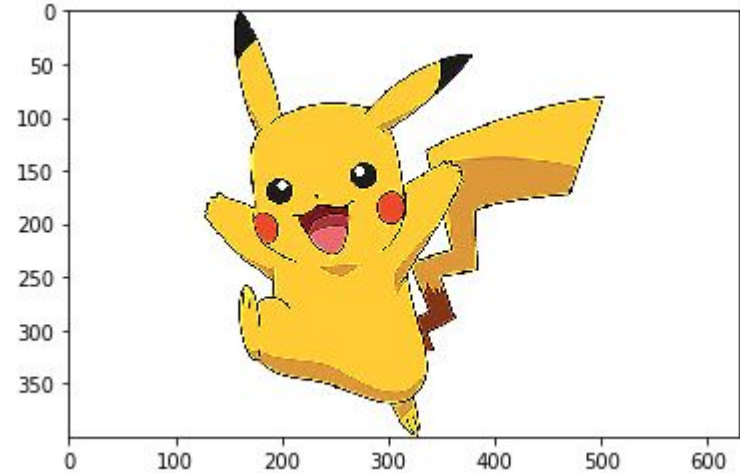


Image

```
print(kernel)
```


```
[[ 0 -1  0]  
 [-1  5 -1]  
 [ 0 -1  0]]
```

Sharpen



**Convolved
Features**

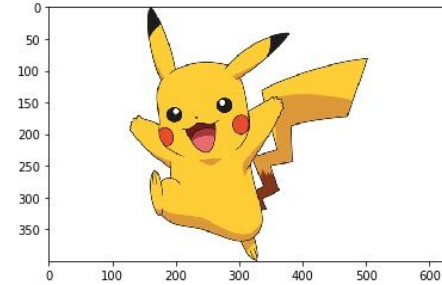
Filter comes from “Image Processing”



Image



Identity



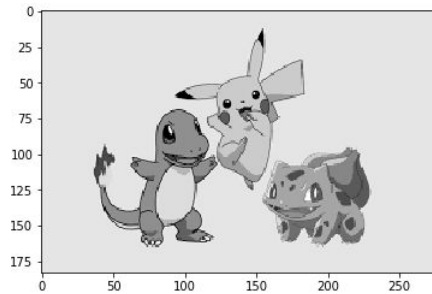
**Convolved
Features**

Where are these filters from?

- Filters, in nature, are model parameters, which can be **learned** by Gradient Descent Algorithms .
- These filters weights are firstly randomly initialized, and then updated during training process.
- End-to-End optimization: Gradients computed by backpropagation.
- More details:
<https://towardsdatascience.com/training-a-convolutional-neural-network-from-scratch-2235c2a25754>

Non-linear Activation

- Filter operation is dot product (linear computation).
- In deep learning, we need to have non-linear transformations.
- Add non-linear activation



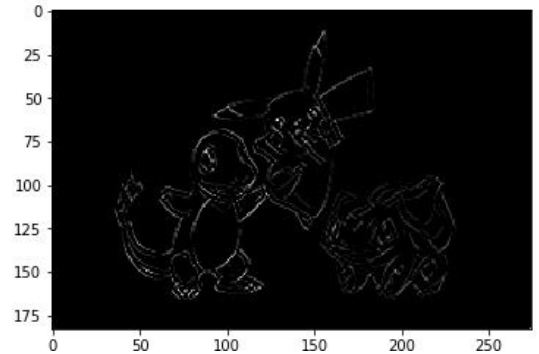
Image

```
print(kernel)
```

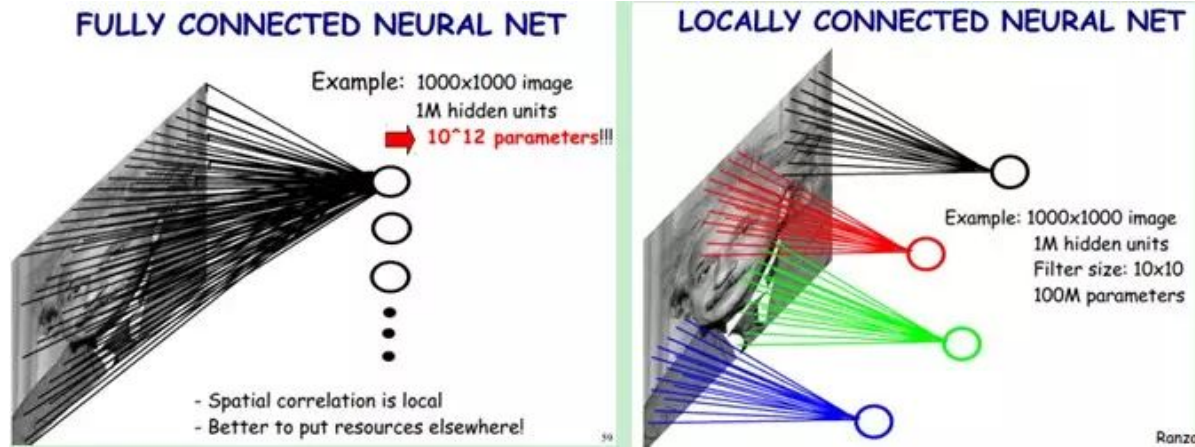
```
[[ 1  0 -1]  
 [ 0  0  0]  
 [-1  0  1]]
```



non-linear



Locally Connected

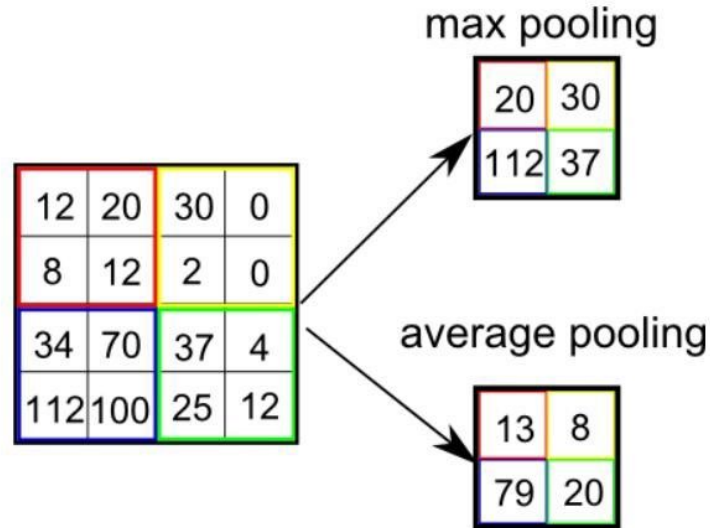


<https://cv-tricks.com/cnn/understand-resnet-alexnet-vgg-inception/>

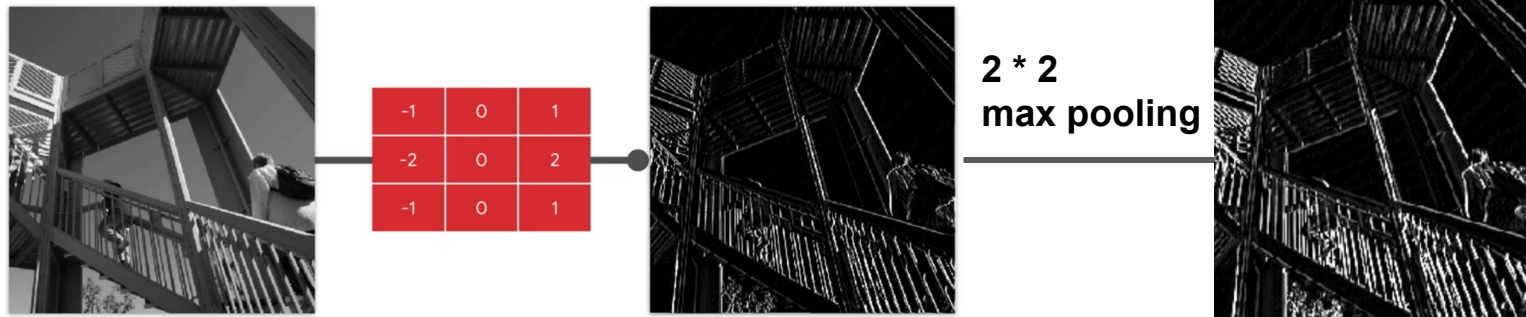
Pooling Operation

- Pooling Size: the box size. Here is $2 * 2$
- Stride Size: how much pixel the window move
- Reduce the dimensionality

What is stride size here ?




Filter then Pool



1. The size is **one quarter** the original size
2. The **vertical line** features are **enhanced**.

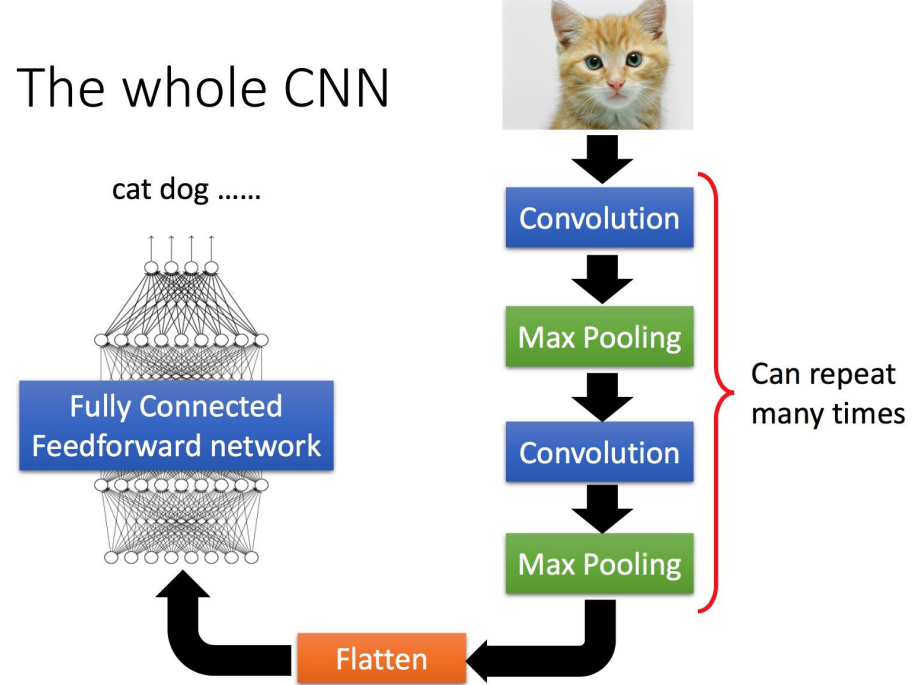
Conv-Pool



CNN Can be Deep

- Convolution-Pooling can be followed by another Convolution-Pooling
- At the end, after flatten operation, fully connected layers are used to map the outputs.

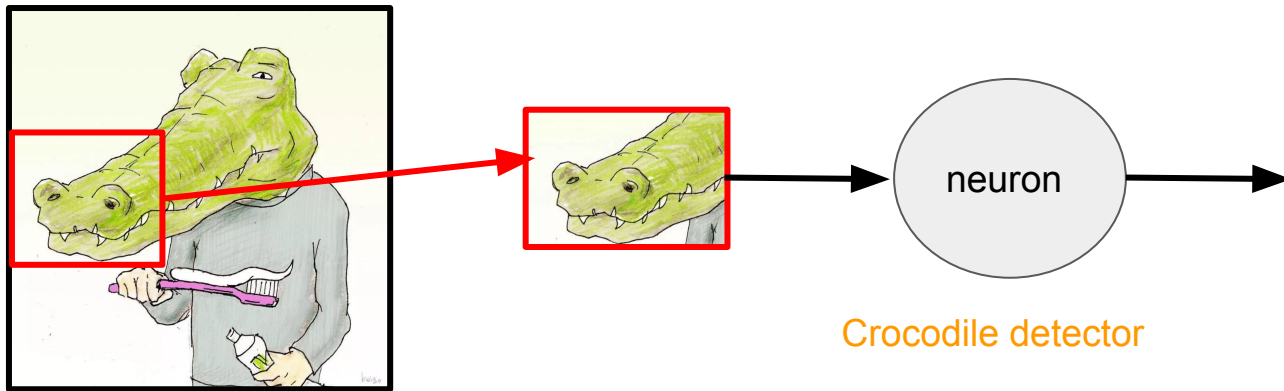
The whole CNN



Why CNN is Suitable for Images

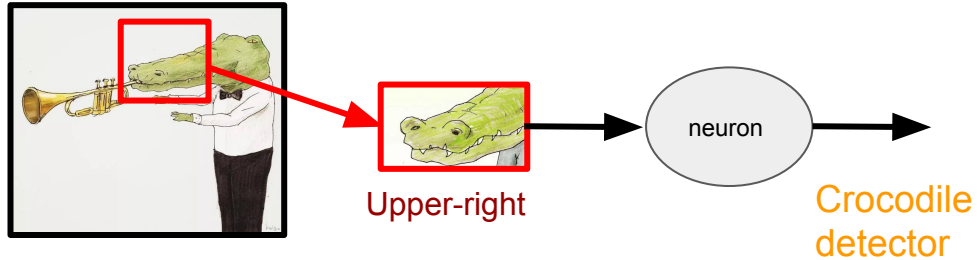
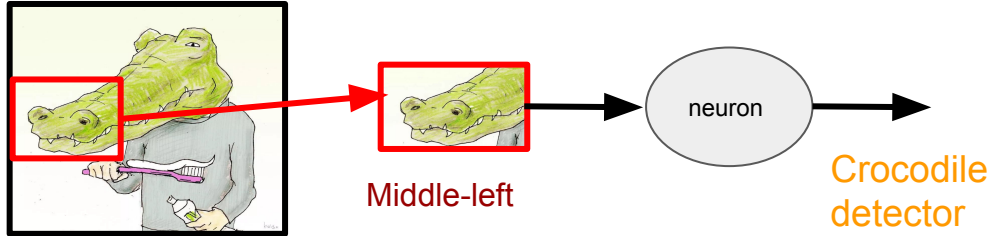
Local Features Matter

- Discriminative patterns are much smaller than the whole image
- A neuron does not have to see the whole image
- Less parameters required



Location Insensitive

- The same patterns appear in different regions
- A neuron should be location insensitive.



Subsampling Works

- Subsampling the pixels will not change the object
- We can subsample the pixels to make images smaller -> less parameters required

Crocodile



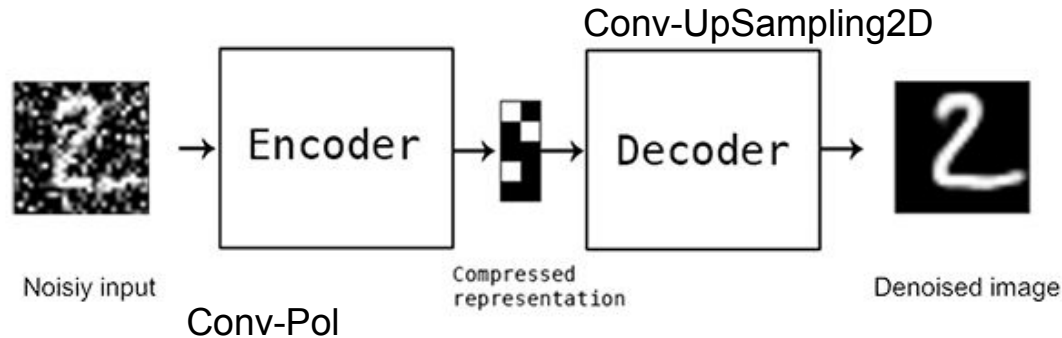
subsampling

Crocodile



Applications

- Image Recognition
- Object Detection
- Image Denoising



<https://blog.keras.io/building-autoencoders-in-keras.html>

<https://www.kaggle.com/michalbrezk/denoise-images-using-autoencoders-tf-keras>

Limitations of CNN

CNN is different human vision

- CNN can handle translations. But they can not cope with the effects of **changing viewpoints such as rotation and scaling**
- Human is able to generalize knowledge.

neatly positioned

ImageNet

Chairs



Real world

ObjectNet

Chairs by rotation



Chairs by background



Chairs by viewpoint



Teapots



T-shirts

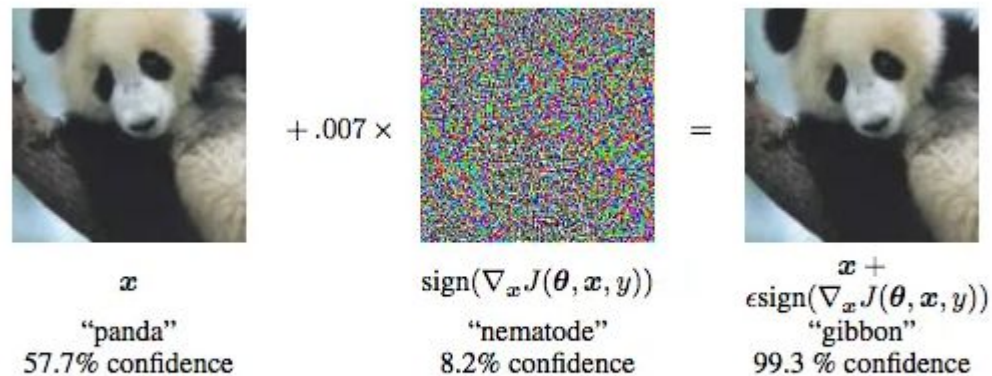


CNN is different human vision

- CNN may get confused by seeing this bizarre teapot, since they can not understand images in terms of objects and their parts.
- Human is able to decompose an object into parts and then we can understand its nature.



CNN is different human vision


$$\begin{array}{ccc} \begin{array}{c} x \\ \text{"panda"} \\ 57.7\% \text{ confidence} \end{array} & + .007 \times \begin{array}{c} \text{sign}(\nabla_x J(\theta, x, y)) \\ \text{"nematode"} \\ 8.2\% \text{ confidence} \end{array} & = \begin{array}{c} \begin{array}{c} x + \\ \epsilon \text{sign}(\nabla_x J(\theta, x, y)) \end{array} \\ \text{"gibbon"} \\ 99.3\% \text{ confidence} \end{array} \end{array}$$

Adversarial examples can cause neural networks to misclassify images while appearing unchanged to the human eye

CNN for Structured Data

Default of Credit Card Clients Dataset

- Static Features
- Dynamic Features

Task: Predict the probability of credit default based on credit card owner's payment status, balance and payment history (for the past 6 months from the predicted period)

Content

There are 25 variables:

- ID: ID of each client
- LIMIT_BAL: Amount of given credit in NT dollars (includes individual and family/supplementary credit)
- SEX: Gender (1=male, 2=female)
- EDUCATION: (1=graduate school, 2=university, 3=high school, 4=others, 5=unknown, 6=unknown)
- MARRIAGE: Marital status (1=married, 2=single, 3=others)
- AGE: Age in years
- PAY_0: Repayment status in September, 2005 (-1=pay duly, 1=payment delay for one month, 2=payment delay for two months, ..., 8=payment delay for eight months, 9=payment delay for nine months and above)
- PAY_2: Repayment status in August, 2005 (scale same as above)
- PAY_3: Repayment status in July, 2005 (scale same as above)
- PAY_4: Repayment status in June, 2005 (scale same as above)
- PAY_5: Repayment status in May, 2005 (scale same as above)
- PAY_6: Repayment status in April, 2005 (scale same as above)
- BILL_AMT1: Amount of bill statement in September, 2005 (NT dollar)
- BILL_AMT2: Amount of bill statement in August, 2005 (NT dollar)
- BILL_AMT3: Amount of bill statement in July, 2005 (NT dollar)
- BILL_AMT4: Amount of bill statement in June, 2005 (NT dollar)
- BILL_AMT5: Amount of bill statement in May, 2005 (NT dollar)
- BILL_AMT6: Amount of bill statement in April, 2005 (NT dollar)
- PAY_AMT1: Amount of previous payment in September, 2005 (NT dollar)
- PAY_AMT2: Amount of previous payment in August, 2005 (NT dollar)
- PAY_AMT3: Amount of previous payment in July, 2005 (NT dollar)
- PAY_AMT4: Amount of previous payment in June, 2005 (NT dollar)
- PAY_AMT5: Amount of previous payment in May, 2005 (NT dollar)
- PAY_AMT6: Amount of previous payment in April, 2005 (NT dollar)
- default.payment.next.month: Default payment (1=yes, 0=no)

<https://www.kaggle.com/uciml/default-of-credit-card-clients-dataset>

Feature Engineering

- Extract as much information as possible from the available datasets, especially dynamic features.
- Given the past 6 months bill payments (a sequence of 6 numbers):
 - The averaged bill payment
 - The difference between two consecutive payments
 -

Content

There are 25 variables:

- ID: ID of each client
- LIMIT_BAL: Amount of given credit in NT dollars (includes individual and family/supplementary credit)
- SEX: Gender (1=male, 2=female)
- EDUCATION: (1=graduate school, 2=university, 3=high school, 4=others, 5=unknown, 6=unknown)
- MARRIAGE: Marital status (1=married, 2=single, 3=others)
- AGE: Age in years
- PAY_0: Repayment status in September, 2005 (-1=pay duly, 1=payment delay for one month, 2=payment delay for two months, ... 8=payment delay for eight months, 9=payment delay for nine months and above)
- PAY_2: Repayment status in August, 2005 (scale same as above)
- PAY_3: Repayment status in July, 2005 (scale same as above)
- PAY_4: Repayment status in June, 2005 (scale same as above)
- PAY_5: Repayment status in May, 2005 (scale same as above)
- PAY_6: Repayment status in April, 2005 (scale same as above)
- BILL_AMT1: Amount of bill statement in September, 2005 (NT dollar)
- BILL_AMT2: Amount of bill statement in August, 2005 (NT dollar)
- BILL_AMT3: Amount of bill statement in July, 2005 (NT dollar)
- BILL_AMT4: Amount of bill statement in June, 2005 (NT dollar)
- BILL_AMT5: Amount of bill statement in May, 2005 (NT dollar)
- BILL_AMT6: Amount of bill statement in April, 2005 (NT dollar)
- PAY_AMT1: Amount of previous payment in September, 2005 (NT dollar)
- PAY_AMT2: Amount of previous payment in August, 2005 (NT dollar)
- PAY_AMT3: Amount of previous payment in July, 2005 (NT dollar)
- PAY_AMT4: Amount of previous payment in June, 2005 (NT dollar)
- PAY_AMT5: Amount of previous payment in May, 2005 (NT dollar)
- PAY_AMT6: Amount of previous payment in April, 2005 (NT dollar)
- default.payment.next.month: Default payment (1=yes, 0=no)

Design of those hand-crafted features is challenging, time-consuming, requires domain knowledge.

Representation of data in CNN format

Shape: 1 by 18

PAY_0
PAY_2
PAY_3
PAY_4
PAY_5
PAY_6
BILL_AMT1
BILL_AMT2
BILL_AMT3
BILL_AMT4
BILL_AMT5
BILL_AMT6
PAY_AMT1
PAY_AMT2
PAY_AMT3
PAY_AMT4
PAY_AMT5
PAY_AMT6


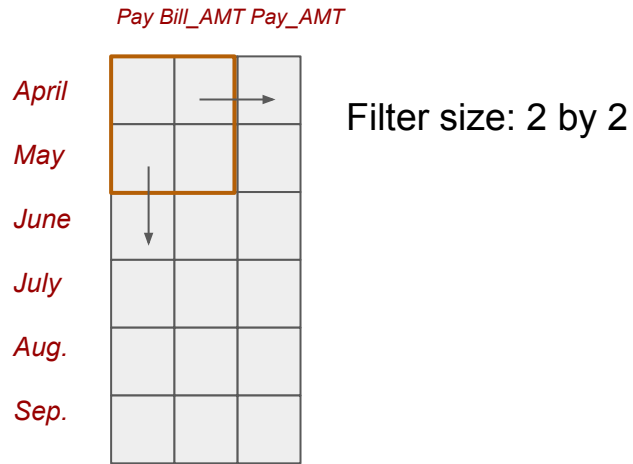


Shape: 1 by 6 by 3

	Pay	Bill_AMT	Pay_AMT
April			
May			
June			
July			
Aug.			
Sep.			

CNN can be easily applied to extract local patterns

Convolution Operation



Which structure is better?

Multiple Channels


- In computer vision, CNN is applied on R-G-B channels
- In this application, different types of credit cards or mortgage of a certain customer can be regarded as different channels



For each customer,
the data shape: 1 by 6 by 3 by **3**

Incorporating Static Features

- Multi-input deep learning is able to combine static and dynamic features for prediction.
- This architecture connects parts of the inputs directly to the output layer.



**Can CNN classify digimon
and pokemon?**

Case Study



<https://medium.com/@DataStevenson/teaching-a-computer-to-classify-anime-8c77bc89b881>

Task Definition

Training Data



Digimon



Pokemon

Testing Data



Digimon or Pokemon?

Build CNN Model

```
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D
from keras.layers import Activation, Dropout, Flatten, Dense

model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(150, 150, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten()) # this converts our 3D feature maps to 1D feature vectors
model.add(Dense(64))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid', name='preds'))

model.compile(loss='binary_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])
```

```
Epoch 1/3
8/8 [=====] - 12s 2s/step - loss: 2.7443 - accuracy: 0.7675 - val_
loss: 0.0834 - val_accuracy: 0.9922
Epoch 2/3
8/8 [=====] - 12s 2s/step - loss: 0.0560 - accuracy: 0.9835 - val_
loss: 0.0692 - val_accuracy: 0.9961
Epoch 3/3
8/8 [=====] - 12s 1s/step - loss: 0.0559 - accuracy: 0.9856 - val_
loss: 0.0684 - val_accuracy: 0.9961
```

Only after three epochs, the testing/val accuracy was easily over 99%. **Amazing!**