

Word Embeddings

Representation Learning for Words

Zhao Rui

Vector Semantic

Word Representation

- How to represent word in a vector space

apple [0 0 0 0 0 **1** 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0]

orange [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 **1** 0 0 0 0 ... 0 0 0 0 0 0]


car [0 0 0 0 0 0 0 0 **1** 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0]

One-hot Vector

- Can we use the above one-hot vector for words?
- Can not capture the semantics of the corresponding words

Distributed Representation

- Words should be encoded into a **low-dimensional** and **dense** vector



Vector Semantics

- Words are characterized by the words that occur with them.
- Words are close to each in the vector space if they are semantically closer to each other.
- It is also called **distributional semantics**.

Motivations

- “**You shall know a word by the company it keeps**” - by Firth (1957)
- Example from Nida (1975); Lin(1998); Jurafsky (2015)

What is Tesgüino?

A bottle of Tesgüino is on the table.

Everybody likes tesgüino

Tesgüino makes you drunk

We make Tesgüino out of corn

- From context words, the meaning behind the word can be inferred as:
 - An alcoholic beverage like beer.



Distributional Semantics

- Words are represented by their context.
- Two words are similar if they have similar word contexts.

I eat an apple every day.

A diagram illustrating the context of the word 'apple' in the sentence 'I eat an apple every day.' Two curved arrows originate from the word 'apple' and point to the words 'eat' and 'every', indicating that these words are part of the context of 'apple'.

I eat an orange every day.

A diagram illustrating the context of the word 'orange' in the sentence 'I eat an orange every day.' Two curved arrows originate from the word 'orange' and point to the words 'eat' and 'every', indicating that these words are part of the context of 'orange'.

Context: Nearby Words


I like driving my car to work.

A diagram illustrating the context of the word 'car' in the sentence 'I like driving my car to work.' Three curved arrows originate from the word 'car' and point to the words 'driving', 'my', and 'to', indicating that these words are part of the context of 'car'.

Bag-of-words

- We represent how often a word occurs in a document
 - Sklearn countvectorizer
 - It is called document-term matrix
- If we look at the column instead of the row

I eat an apple every day
I eat an orange every day
I like driving my car to work



	an	apple	car	day	driving	eat	every	like	my	orange	to	work
0	1	1	0	1	0	1	1	0	0	0	0	0
1	1	0	0	1	0	1	1	0	0	1	0	0
2	0	0	1	0	1	0	0	1	1	0	1	1

Term-Document Matrix

- Each document is a count vector in a vector space whose dimension is D
 - D is the number of documents
 - The shape of the matrix is $|V| * D$
 - Each row is the vector for the word
 - **Two words are similar if their vectors are similar**

	0	1	2
an	1	1	0
apple	1	0	0
car	0	0	1
day	1	1	0
driving	0	0	1
eat	1	1	0
every	1	1	0
like	0	0	1
my	0	0	1
orange	0	1	0

Word-word matrix

- Document is a kind of “context”. However, it is too abstract.
- Smaller context will be better:
 - **Window of k nearby words**, here k can be 2,3,4,..
- Instead of term-document matrix, we are going to have word-word matrix
 - Each word vector's dimension will be $|V|$
 - The matrix will be the shape of $|V| * |V|$
- To build the word-word matrix:
 - Co-occurrence: For a given corpus, the co-occurrence of a pair of words say w_1 and w_2 is the number of times they have appeared together in a Context Window .
 - Context Window: Context window is specified by a number and the direction (usually set to be left and right).

For example

- For context window: the window size is 2 and the direction is set to be right and left.

I eat an apple **every** day

I eat an orange **every** day

I like driving my car to work



The co-occurrence: (every, an): 2
(every, apple): 1
(every, day): 2

Word-word matrix

- Size will be $V * V$.
- High-dimensional and very sparse
- Symmetry

contexts

	like	an	to	my	driving	apple	orange	work	every	car	I	eat	day
like	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
an	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	2.0	0.0	2.0	2.0	0.0
to	0.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0
my	1.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
driving	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0
apple	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	1.0	1.0
orange	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	1.0	1.0
work	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
every	0.0	2.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	2.0
car	0.0	0.0	1.0	1.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
I	1.0	2.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0
eat	0.0	2.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	2.0	0.0	0.0
day	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	2.0	0.0	0.0	0.0	0.0

words

I eat an apple every day

I eat an orange every day

I like driving my car to work

```
vec_apple = mat[vocab.index('apple')].reshape(1, -1)
vec_orange = mat[vocab.index('orange')].reshape(1, -1)
vec_car = mat[vocab.index('car')].reshape(1, -1)
print('cosine scores between apple and orange vectors')
print(cosine_similarity(vec_apple, vec_orange))
print('cosine scores between apple and car vectors')
print(cosine_similarity(vec_apple, vec_car))
```

```
cosine scores between apple and orange vectors
[[1.]]
cosine scores between apple and car vectors
[[0.]]
```

The size of window

- Under different window size, we will have different word-word matrix

	like	an	to	my	driving	apple	orange	work	every	car	I	eat	day
like	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
an	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	2.0	0.0
to	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0
my	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
driving	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
apple	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
orange	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
work	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
every	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	2.0
car	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
I	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0
eat	0.0	2.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0	0.0
day	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0	0.0	0.0	0.0

Win size=1


	like	an	to	my	driving	apple	orange	work	every	car	I	eat	day
like	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0
an	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	2.0	0.0	2.0	2.0	2.0
to	0.0	0.0	0.0	1.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0
my	1.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	1.0	0.0	0.0
driving	1.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0
apple	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	1.0	1.0
orange	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	1.0	1.0	1.0
work	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
every	0.0	2.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	2.0	2.0
car	1.0	0.0	1.0	1.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
I	1.0	2.0	0.0	1.0	1.0	1.0	1.0	0.0	0.0	0.0	0.0	2.0	0.0
eat	0.0	2.0	0.0	0.0	0.0	1.0	1.0	0.0	2.0	0.0	2.0	0.0	0.0
day	0.0	2.0	0.0	0.0	0.0	1.0	1.0	0.0	2.0	0.0	0.0	0.0	0.0

Win size=3

- From Jurafsky (2015): the size of windows depends on your goals
 - The shorter the windows, the more syntactic the vector (1-3)
 - The longer the windows, the more semantic the representation (4-10)

From Raw Count to PPMI

- Raw word frequency is not a great measure of association between words
 - Very **skewed** distribution. For example, **the** and **of** are very frequent, but may not be the most discriminative
- We can adopt a new measure to replace the counts as the entry in the word-word matrix
 - Positive Pointwise Mutual Information (PPMI)

$$PPMI(w_1, w_2) = \max\left(\log_2 \frac{p(w_1, w_2)}{p(w_1)p(w_2)}, 0\right)$$


Penalize high-frequency words


From Sparse Vectors to Dense Ones

- Matrix decomposition can be applied on the word-word matrix.
- Singular Value Decomposition (SVD) is one of the classic methods.
 - Change the dimensions such that they are orthogonal to each other.
 - The new vector space will keep the first k dimensions that explain the largest amount of variance in the data.


	dim0	dim1	dim2	dim3	dim4
like	-0.563652	1.497387	-0.423378	0.119592	-0.713812
an	-3.590090	-0.398018	-1.152087	1.707900	0.249389
to	-0.225862	1.595626	0.189236	-0.058551	-0.175054
my	-0.609500	2.030081	-0.494240	0.142318	0.746323
driving	-0.590864	1.794945	-0.501255	0.146783	-0.461930
apple	-2.011314	-0.236121	-0.338984	0.457370	-0.054726
orange	-2.011314	-0.236121	-0.338984	0.457370	-0.054726
work	-0.154674	1.260900	-0.017046	0.009616	-0.839717
every	-3.026532	-0.659660	1.762966	0.750429	-0.232352
car	-0.293771	1.874864	0.363469	-0.114090	0.852037
I	-2.620567	0.832339	2.084192	-0.552119	0.073469
eat	-3.081751	-0.211562	-1.373831	-1.499121	-0.042218
day	-2.363795	-0.593157	-0.158542	-1.849308	0.044472

When k is set to be 5 our toy example.

Word Vectors



Word Analogy




Mikolov & Chen et al. 2013
Mikolov & Sutskever et al. 2013

From Sparse Vectors to Dense Ones

- Another approach is prediction based methods instead of matrix methods.
- We would like to build a machine learning model for the task that given certain words, can we predict their context words?
- What is the most powerful supervised prediction model?
 - Neural network
- It is the Wor2Vec model: a neural network based word embedding model.

Neural Network Solution



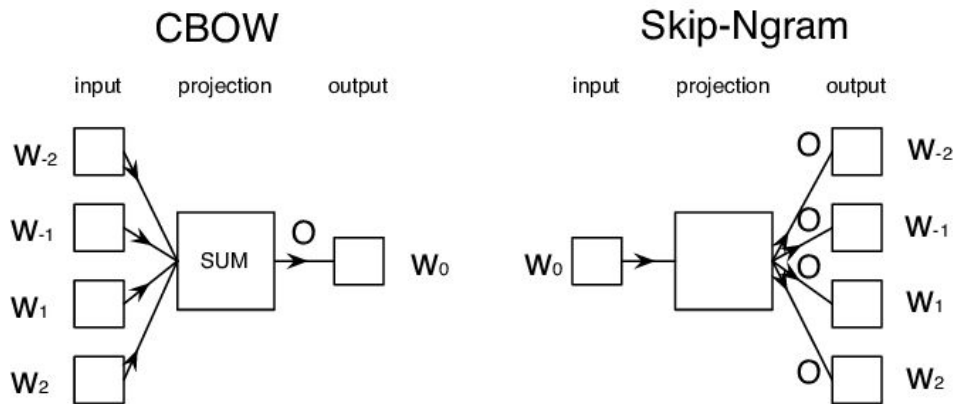
Input word and output words should be sampled from the same context

Another self-supervised learning example

What is Word2Vec?

Word2Vec



- A method of computing vector representation of words developed by Google.
- Open-source version of Word2Vec hosted by Google (in C)
- Train a simple neural network with a single hidden layer to perform word prediction tasks
- Two structures proposed **Continuous Bag of Words (cbow)** vs **skip-gram**:



Word2Vec as BlackBox



input, target



Corpus


Word2Vec Tool

Word Embeddings

A Good Visualization for Word2Vec

<https://ronxin.github.io/wevi/>

Model Architecture




Structure Highlights:

- input layer
 - one-hot vector
- hidden layer
 - linear (identity)
- output layer
 - softmax


Hidden Layer

- There is not activation function here
- 300 neurons are the word vec. dimensions
- This layer is operating as a 'lookup' table
- Input word matrix denoted as **IVec**

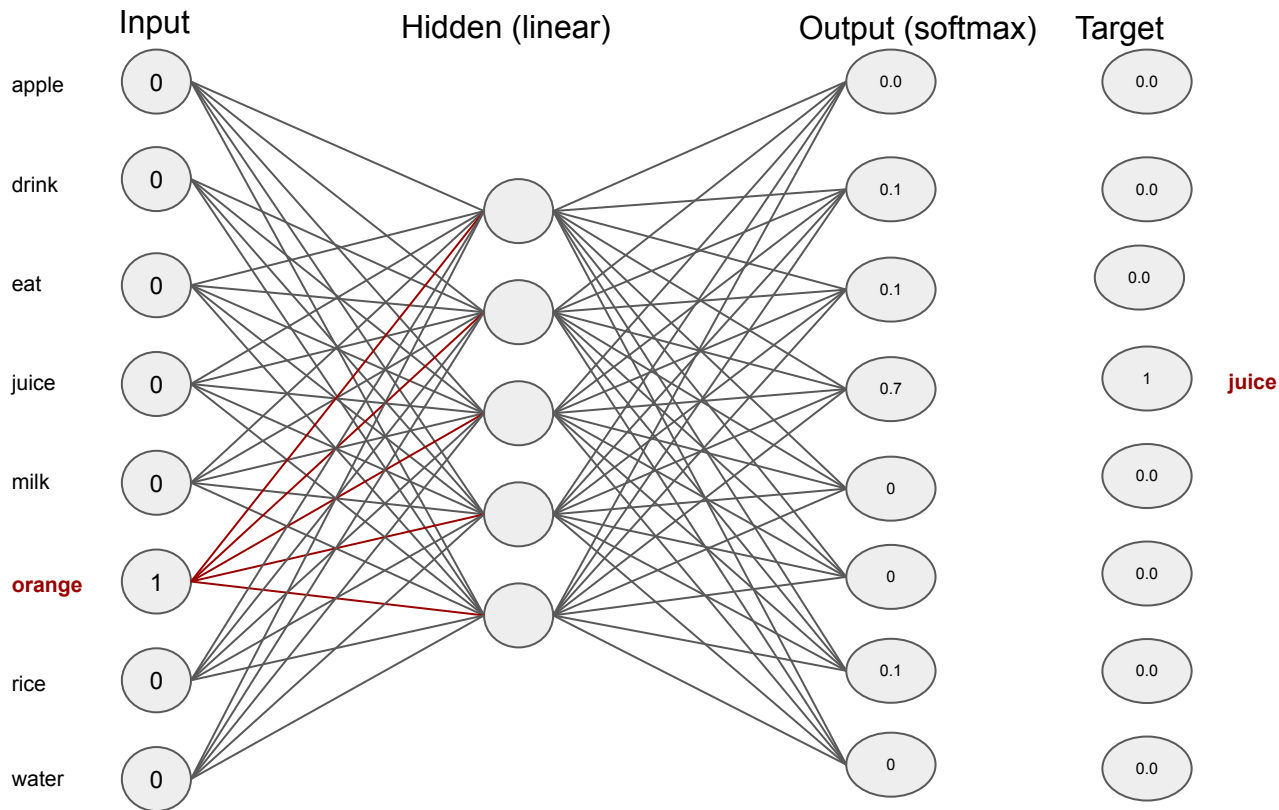


Output Layer

- Softmax classifier
- Output word matrix denoted as **OVec**




Word2Vec Network




Then, we can compute the **loss** and call gradient descent to update model parameters.

Updating Word Vectors



A force-directed graph



Idea behind Word2Vec

- Feature vector assigned to a word will be adjusted if it can not be used for accurate prediction of that word's context.
- Each word's context in the corpus is the teacher sending error signals back to modify the feature vector.
- It means that words with **similar context** will be assigned **similar vectors**!

Distributional Semantics

Input vs Output Word Vectors

- Inputs: semantics encoder from one-hot/word index to semantics
- Outputs: semantics decoder from semantics to probability distributions over words.
- **In most cases, input word vectors are used.** Some have observed that combinations of these two vectors may perform better.

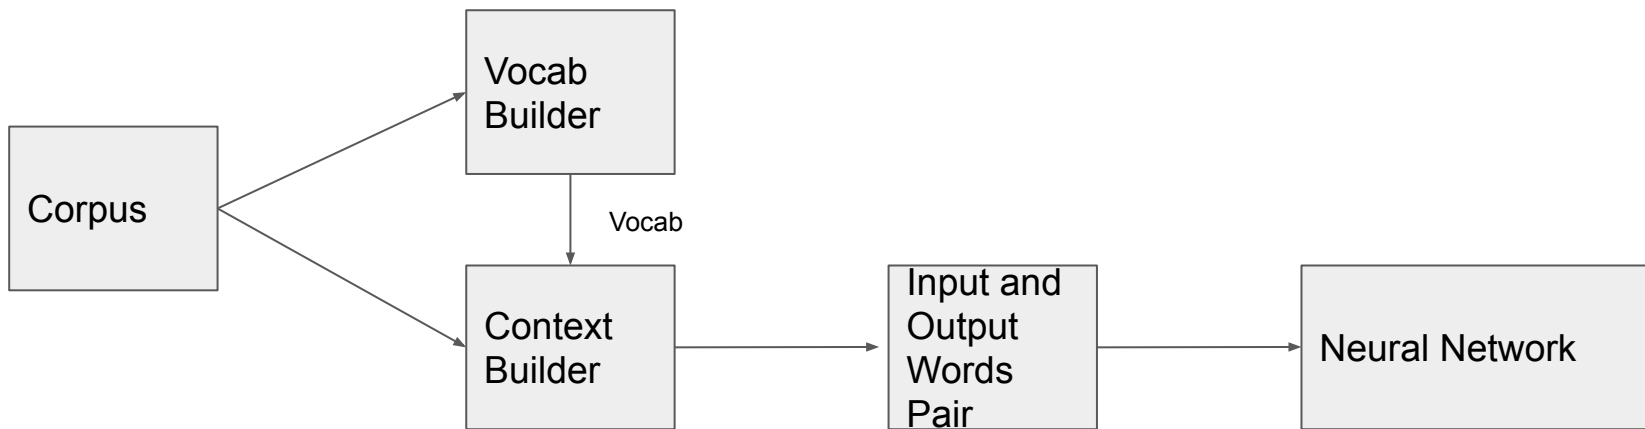
	Vector size	Overall	Semantic	Syntactic
DVRS	300	0.41	0.59	0.26
DVRS	1024	0.43	0.62	0.28
SG	300	0.64	0.69	0.60
SG	1024	0.57	0.60	0.55
Add 300-DVRS, 300-SG	300	0.64	0.72	0.58
Concatenate 300-DVRS, 300-SG	600	0.67	0.74	0.60
Add 1024-DVRS, 1024-SG	1024	0.60	0.66	0.55
Concatenate 1024-DVRS, 1024-SG	2048	0.61	0.68	0.55
Concatenate DVRS-1024, SG-300	1324	0.66	0.73	0.60
Oracle DVRS-1024, SG-300	1024/300	0.70	0.79	0.62

Garten, 2014

Table 2: Performance on word analogy problems with vectors trained against the first 10⁹ bytes of Wikipedia.


Input and Output Words

- How to select them from corpus
- **Skip-gram** and **CBoW** differs here.



Skip-Gram


- Task Definition: given a specific word, predict its nearby word (probability output)
- Model input: source word, Model output: nearby word
- Input is one word, output is one word
- The output can be interpreted as prob. scores, which are regarded as that how likely it is find each vocabulary word can be nearby your input word.



CBoW


- Task Definition: given context, predict its target word
- Model input: context (several words), Model output: center word
- Input is several words, output is one word
- Core Trick: **average** these context vectors for prob score computing

give a **talk** at the



Input x
(give,a,at,the)

Target y
talk



Skip-Gram Vs CBoW

- **CBoW: learning to predict the word by the context**
- **Skip-gram: learning to predict the context by the center word**
- CBoW: several times faster to train the skip-gram
- Skip-gram: works well with small amount of the training data, represents well even rare words or phrases.

Context Selection


- In count-based or predict-based methods, context has a large effect.
- Small context window: more syntax-based embeddings
- Large context window: more semantics-based, topical embeddings
- **Engineering practice: window size is randomly sampled between 1 and maximum window size**

Huge Number of Parameters

- Vocab size is huge
- The Sum of operation in softmax layer is very expensive, i.e., $O(v)$.

$$e^{(\text{IVec}[\text{eat}] * \text{OVec}[\text{apple}])} + e^{(\text{IVec}[\text{eat}] * \text{OVec}[\text{juice}])} + e^{(\text{IVec}[\text{eat}] * \text{OVec}[\text{drink}])} + e^{(\text{IVec}[\text{eat}] * \text{OVec}[\text{other vocab words}])}$$

- Two solutions: **Hierarchical softmax** and **negative sampling**



Document Embeddings


Word Embeddings

- Glove, Fasttext and other open-source nlp methods can learn dense and low-dimensional vectors for words
- Starting from word embeddings, we can learn vectors for higher-level natural language units such as sentences and documents.

Sequence of Words


- Each sentence or document can be regarded as a sequence of vectors.

I hate this movie



4 by d

This is my favorite movie.



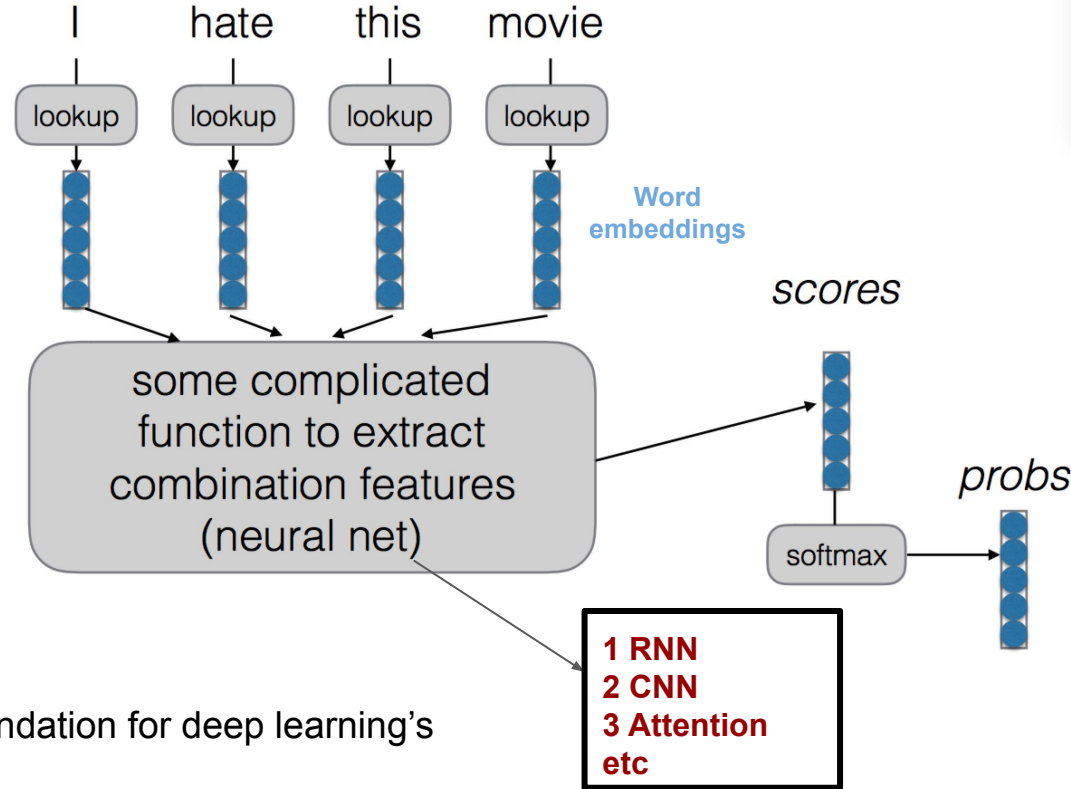
5 by d

- The shape of matrix depends on the length of sequence. However, the majority of ML systems need fixed-length feature vectors.
- One simple solution: average the sequence of vectors, just like bag-of-words.

Complex Semantic

1. **Input Text:** a sequence of words;
2. **Through Word Embedding Look-up:** a sequence of word vectors;
3. Neural networks is applied upon the vector sequences to learn semantic **composition** for final prediction;

→ Human understand the word meaning firstly, then get the whole sentence meaning by composing these words' meaning together.



Word Embeddings is the foundation for deep learning's applications on NLP

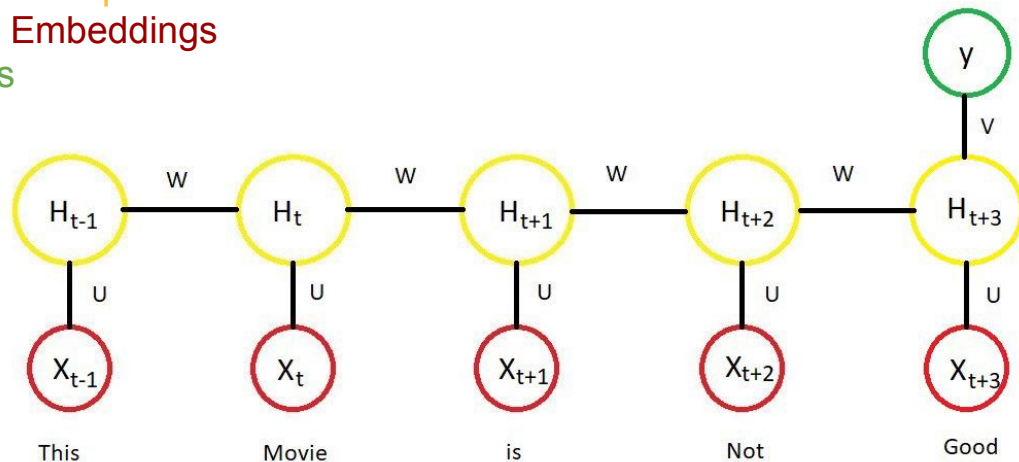
Recurrent Neural Network for NLP

U, W, V: RNN's parameters

H: Hidden Outputs

X: Word Embeddings


y: Labels



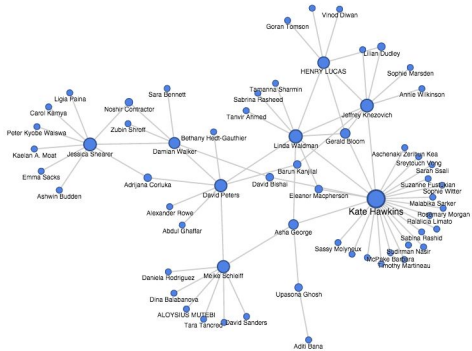
Graph Embedding

Graph Data

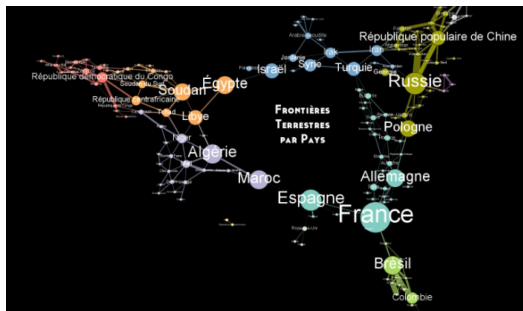
- Graph is an ordered pair $G=(V, E)$.
- V is the set of nodes
- E is the collection of paris of nodes which is called edges



Graphs Are Everywhere




Social Network Analysis



Logistics and Transportation



Recommendation System



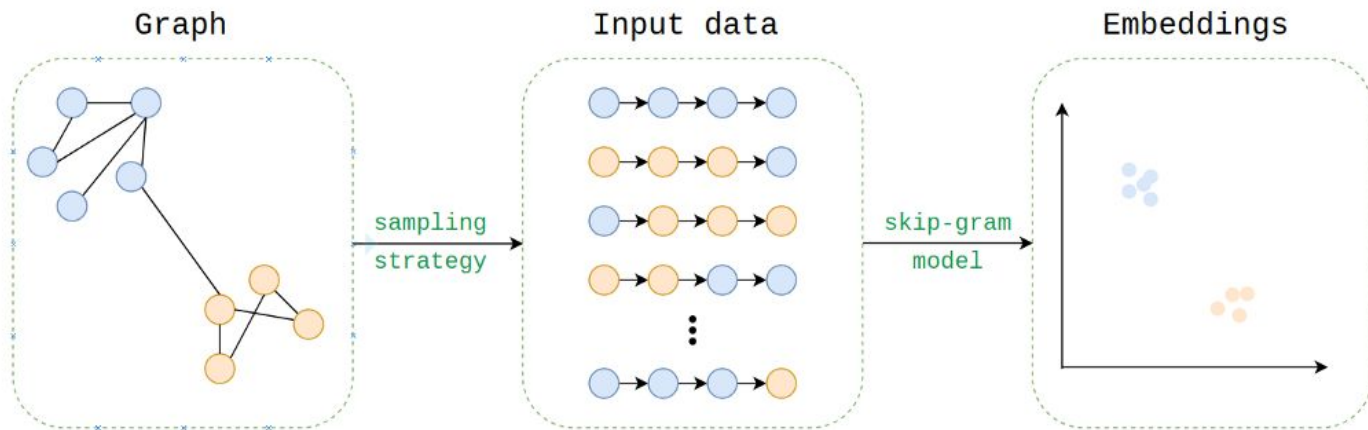
Fraud Detection

Graph Data

- Based on your tasks, define your nodes and Edges
- Apply graph mining algorithms:
 - Graph Pattern Mining
 - Graph Classification
 - Graph Compression
 - Graph Clustering
 - Etc

Embedding for Graph Data

- Embeddings can be extended beyond NLP domain
- Embeddings can be learned for any nodes in a graph



- Node can be items, web page and so on in user clicked stream data
- Embeddings can be learned for any group of discrete and co-occurring states.