

Generative Deep Learning

Agenda

1. Generative Models
2. Generative Adversarial Network

Generative Models

Generative Models

- Given training data, generate new samples from same distribution



Training data $\sim p_{\text{data}}(\mathbf{x})$




Generated samples $\sim p_{\text{model}}(\mathbf{x})$


- Want to learn $p_{\text{model}}(\mathbf{x})$ similar to $p_{\text{data}}(\mathbf{x})$
- Address density estimation, a core problem in ? learning

Generative Models

- Realistic samples for artwork, super-resolution, colorization, etc



- Generate data samples that can be used for simulation and planning (reinforcement learning applications)



Taxonomy of Generative Models

*Explicitly define and solve
for $p_{\text{model}}(x)$*




Figure copyright and adapted from Ian Goodfellow, Tutorial on Generative Adversarial Networks, 2017.

Generative Adversarial Network

LeCun's Comment

What are some recent and potentially upcoming breakthroughs in deep learning



Yann LeCun, Director of AI Research at Facebook and Professor at NYU
Answered Jul 29, 2016 · Upvoted by Joaquin Quiñero Candela, studied Machine Learning and Thamme Gowda, M.S. Computer Science, University of Southern California (2017)




There are many interesting recent development in deep learning, probably too many for me to describe them all here. But there are a few ideas that caught my attention enough for me to get personally involved in research projects.

The most important one, in my opinion, is adversarial training (also called GAN for Generative Adversarial Networks). This is an idea that was originally proposed by Ian Goodfellow when he was a student with Yoshua Bengio at the University of Montreal (he since moved to Google Brain and recently to OpenAI).

This, and the variations that are now being proposed is the most interesting idea in the last 10 years in ML, in my opinion.

Generative Adversarial Networks

- GANs: do not try to find density function, but only want to sample from distribution
- Two sub-modules:
 - Generator: sample from a simple distribution (random gaussian). Learn transformation to training distribution.



- Discriminator: detect the sample whether it is from real distribution



Generator vs Discriminator

- Generator: try to fool the discriminator by generating real-looking images
- Discriminator: try to distinguish between real and fake images

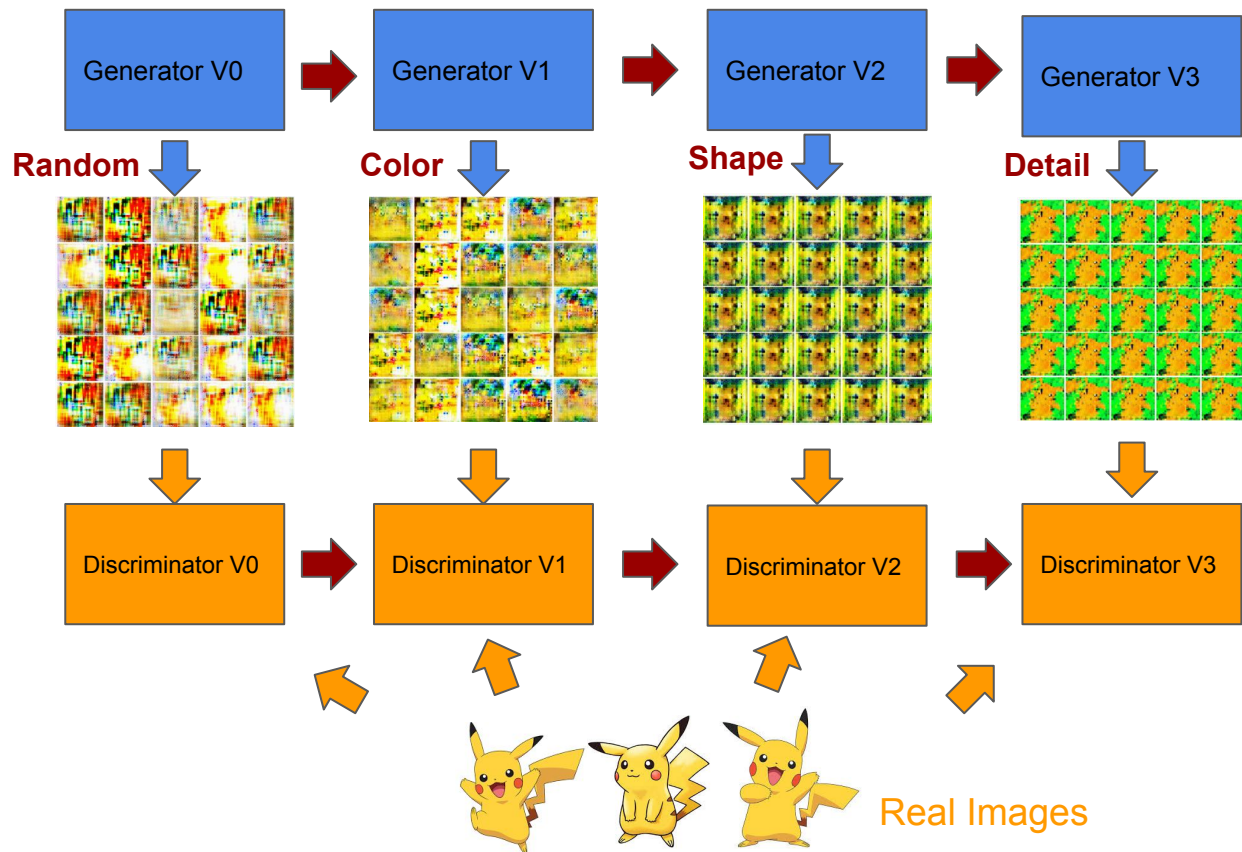
Frenemy

- Generator: try to fool the discriminator by generating real-looking images
- Discriminator: try to distinguish between real and fake images

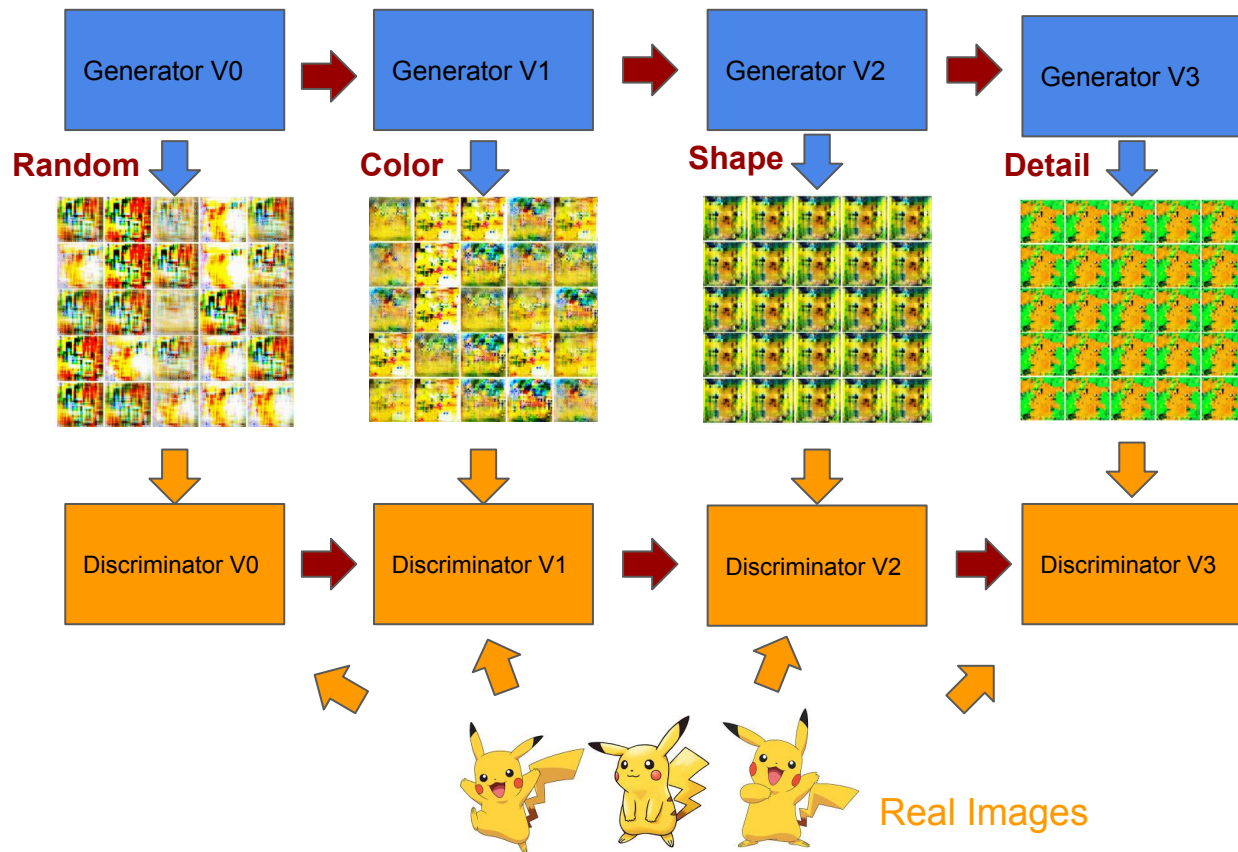


Frenemy
Love-hate relationship

Ideas behind GAN




Ideas behind GAN



Generator and Discriminator are **transformation**, i.e., **neural networks**

Algorithms for GAN


- Get the real data collection:



- Randomly initialize generator network G and discriminant network D
- During training, iteratively train G and D

Algorithms for GAN


- Step 1: Fix Generator G, Train discriminant D



Discriminator are updated to assign high scores to real objects and low scores to generated objects


Algorithms for GAN

- Step 2: Update Generator G, Fix discriminant D




Generator learn to fool the discriminator


Train GAN over mnist dataset




Epoch 0




Epoch 500




Epoch 1000




Epoch 1500



Final Output




Final Output



It is very hard to train GAN

- Non-convergence: the model parameters oscillate and never converge
- Model collapse: the generator collapses which produces limited modes of samples



Good
GAN

Bad
GAN

- Diminished gradient: the discriminator gets too successful that the generator gradient vanishes and learns nothing.

Game Theory and GAN

- GAN is the minimax/zero-sum non-cooperative game
- GAN's minimax equation as:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_r(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

- D's actions are to maximize them and G wants to minimize its actions
- In game theory, GAN model converges when the D and G reach a Nash Equilibrium

Nash equilibrium

- In minimax game, both sides want to beat the others
- Nash equilibrium is that when one player will not change its action regardless of what the opponent may do
- Consider two players A and B which control the value of x and y, Player A wants to maximize the value xy while B wants to minimize it

$$\min_B \max_A V(A, B) = xy$$

- What are the values of x and y for the nash equilibrium?


Gradient Descent for Nash equilibrium

- Update parameter x and y following the gradient of the value function

$$\uparrow \Delta x = \alpha \frac{\partial(xy)}{\partial x}$$

$$\downarrow \Delta y = -\alpha \frac{\partial(xy)}{\partial y}$$

- During training iterations, start with the initial guess of $x=2$ and $y=2$



Can not Converge i.e. $x=y=0$



2014



2015



2016



2017



2018

GAN progress on face generation from Ian Goodfellow